

Company & Product Overview

What Sarvam AI actually sells, who pays for it, and why a company competing on Indian-language depth and data sovereignty is a genuinely interesting product problem.

COMPANY Sarvam AI, Bengaluru	FOUNDED August 2023	CATEGORY Full-stack sovereign AI	BUYER TYPES Developer · Enterprise · Government	AS OF August 2026
---------------------------------	------------------------	-------------------------------------	--	----------------------

1.1 What is Sarvam AI?

Sarvam AI is a Bengaluru-based startup building **India's full-stack sovereign AI platform** — foundation models, speech and voice AI, translation, document intelligence, and developer APIs, purpose-built for Indian languages, code-mixed text (Hinglish, romanized Hindi), and India-specific deployment needs such as data residency and regulated-industry compliance.

Unlike a single-product company, Sarvam operates across three layers. This matters for a PM because each layer has a different customer, a different sales motion, and a different definition of success.

LAYER 03 Applications & GTM	INDUS — CONSUMER ASSISTANT SARVAM KAZE — AI GLASSES	SAMVAAD — ENTERPRISE VOICE AGENTS FORWARD-DEPLOYED ENGINEERING
LAYER 02 Foundation Models	SARVAM-30B / 105B — REASONING LLMS SARVAM TRANSLATE	SAARAS — ASR BULBUL — TTS SARVAM VISION — OCR
LAYER 01 Infrastructure	INDIAAI MISSION GPUS SOVEREIGN / AIR-GAPPED DEPLOYMENT	INFERENCE SERVING

1.2 Who is the customer?

Sarvam sells to three distinct buyer types — itself a strategic complexity worth naming, because a single roadmap has to serve all three.

Developers SELF-SERVE	Enterprises SALES-LED	Government PROCUREMENT
Build voice, translation, or chat features on pay-as-you-go APIs priced in rupees. Low ARPU, high volume, product-led.	Banking, insurance, telecom, and consumer tech buying multilingual voice agents or document AI via contracts plus forward-deployed engineers.	Sovereign, air-gapped deployments for schemes and citizen services reaching hundreds of millions of people.

1.3 What problem does it solve?

- Global frontier models (GPT, Claude, Gemini) are optimised for English and Western languages first. **Indian-language quality, code-mixed input handling, and voice-first interaction for non-typing users are secondary**

for them.

- India has roughly **800 million people who are more comfortable speaking than typing**, across 22 scheduled languages — a market global providers under-serve.
- Regulated Indian institutions need **data sovereignty**: models and inference that never leave Indian soil. Foreign API providers structurally cannot offer this.

CORE PROBLEM STATEMENT

Indian businesses and citizens need AI that understands India's languages, scripts, and voice-first usage patterns — and regulated institutions need that AI to run entirely within Indian jurisdiction.

1.4 Why does the product exist?

Founders **Vivek Raghavan** and **Pratyush Kumar** came from AI4Bharat at IIT Madras, where they spent years building multilingual speech and text datasets — but found that academic research wasn't reaching end users, such as a farmer trying to understand a government scheme. Sarvam exists to convert that research into a commercial, production-grade product layer, backed by IndiaAI Mission compute subsidies and a national push for token sovereignty.

1.5 Business model at a glance

REVENUE STREAM	DESCRIPTION
Pay-as-you-go API	₹ per token, per audio-minute, per character, and per page across chat, STT, TTS, translation, and document parsing
Enterprise contracts	Custom pricing, volume discounts, dedicated support, on-prem or air-gapped deployment for banks, insurers, and government
Consumer app (Indus)	Currently free / beta; monetisation not yet public — likely future freemium, subscription, or B2B2C distribution
Hardware (Sarvam Kaze)	Smart glasses — a potential new hardware revenue line, still early
Government programmes	Possibly grant or mission-funded deployments tied to IndiaAI Mission rather than direct revenue

1.6 Competitive set

CATEGORY	COMPETITORS
Global frontier LLM providers	OpenAI, Anthropic, Google Gemini
Voice & speech specialists	SoundHound, Sanas, AssemblyAI, ElevenLabs, Google Cloud Speech
Indian-language / sovereign peers	Krutrim (Ola), other IndiaAI Mission-backed model builders
Hyperscalers (infra + AI bundled)	AWS, Google Cloud, Microsoft Azure

1.7 Key differentiation

Trained natively, not fine-tuned

Built from scratch on Indian languages rather than adapted from Llama, Mistral, or Qwen — meaningfully better on code-mixed and romanized Indian text.

Widest production ASR coverage in India

22 languages via Saaras, ahead of every global and domestic alternative on breadth.

Sovereignty as a structural moat

Air-gapped, India-hosted deployment for regulated sectors — something foreign providers cannot easily replicate.

INR pricing, forward-deployed delivery

Transparent rupee pay-per-use removes FX and GST friction; embedded engineers give large customers faster iteration than a typical API vendor.

1.8 The honest limitation

BALANCED READ

Sarvam's own materials and third-party reviews note it does **not** aim for universal frontier dominance. It competes on Indian-language depth and sovereignty, not on beating Claude, GPT, or Gemini at general-purpose global tasks. A credible teardown holds both truths at once: strong niche differentiation, alongside real technical and ecosystem gaps in documentation maturity and model breadth versus global leaders.

CARRIES FORWARD TO

The three-buyer structure drives the funnel split in **Section 5**, the bifurcated unit economics in **Section 13**, and the decision in **Section 9** to tear down the self-serve developer journey specifically — it is the surface a product manager can move fastest.

Sources: sarvam.ai, docs.sarvam.ai, Sarvam Series B announcement, Tracxn company profile. Verify current details on the official site before quoting externally.