

Jobs To Be Done

What each persona is actually hiring Sarvam to do — expressed as an outcome they are accountable for, not as a product feature they want to use.

FRAMEWORK JTBD (Christensen)	STRUCTURE When... I want to... so that...	JOBS MAPPED 4, one per persona	FEEDS INTO Sections 9, 10 & 12
---------------------------------	--	-----------------------------------	-----------------------------------

The framing used throughout is *"When [situation], I want to [motivation], so that [desired outcome]."* Writing the job this way forces the outcome to be something the buyer is measured on — which is the test any roadmap item in **Section 12** has to pass.

3.1 The four jobs

ROHAN • INDEPENDENT DEVELOPER	
WHEN	I'm building a feature that needs to understand spoken or written Hindi, Tamil, or code-mixed input,
I WANT TO	call a reliable, well-documented API within minutes,
SO THAT	I can ship the feature without training my own speech or language model.

ANANYA • ENTERPRISE BUYER	
WHEN	my call-centre costs are rising and customers in Tier 2 and Tier 3 cities can't be served well in English,
I WANT TO	deploy a multilingual voice-agent platform that meets our data-residency requirements,
SO THAT	I can cut service costs while staying compliant and improving customer reach.

KARTHIK • STARTUP BUILDER	
WHEN	I'm designing a product for users who prefer speaking over typing in their own language,
I WANT TO	build on an Indic-first AI stack with startup-friendly pricing and credits,
SO THAT	I can validate my product with real Bharat users without a large upfront AI bill.

MR. IYER • GOVERNMENT	
WHEN	my department needs to reach hundreds of millions of citizens in their native language without letting citizen data leave Indian soil,
I WANT TO	commission a sovereign, auditable AI deployment,
SO THAT	I can meet policy mandates while actually improving citizen service.

3.2 Job vs. product-catalogue framing

The gap between how a buyer states their job and how Sarvam's surfaces are organised is the root cause of the onboarding friction diagnosed in **Section 9**.

PERSONA	HOW THEY DESCRIBE THE JOB	HOW THE PRODUCT CURRENTLY PRESENTS ITSELF
Rohan	"I need a Hindi voice bot in my app"	Six separate API product pages: chat, Saaras, Bulbul, Translate, Vision, video localisation
Ananya	"I need to cut cost per resolved call and pass a compliance review"	Model capabilities, language counts, deployment options
Karthik	"I need to validate a Bharat product cheaply"	Per-token, per-minute, per-character price tables
Mr. Iyer	"I need an auditable system my department can defend"	Sovereign / air-gapped deployment as a capability claim

3.3 Why JTBD matters here

THE INSIGHT THIS SECTION PRODUCES

None of these jobs is "use an LLM." Each is a business outcome: ship a feature, cut costs, validate a product, meet a policy mandate. That reframing is what justifies the redesign in **Section 10** — Sarvam's onboarding and messaging should map to these jobs, not to model specifications.

What this changes upstream

Documentation and landing pages get organised by use case ("building a voice bot") rather than by product ("Saras ASR"), because that is the language in which the job arrives.

What this changes downstream

Success metrics measure completed jobs — successful tasks, not raw API calls — which is precisely the North Star proposed in **Section 6**.