

Opportunity Matrix

Seven candidate opportunities scored with RICE — including one deliberately low-value item, because a prioritisation framework that ranks everything as important isn't a framework.

FRAMEWORK	FORMULA	OPPORTUNITIES	TOP SCORE
RICE	$(R \times I \times C) \div E$	7	243.0

11.1 Why RICE and not ICE

RICE was chosen over ICE because it forces an explicit **Reach** estimate. That matters here: Sarvam's audience splits between a large self-serve developer base and a small number of very high-value enterprise deals, and ICE alone would flatten that distinction — making a change affecting eight thousand developers look identical to one affecting eight enterprise accounts.

Scale used: Reach, Impact, and Confidence are scored 1–10 where 10 is highest. Effort is scored 1–10 where 10 is the most effort, roughly a full quarter or more of work.

11.2 Scored opportunities

OPPORTUNITY	REACH	IMPACT	CONFIDENCE	EFFORT	RICE SCORE
Live credit / cost estimator at signup	9	6	9	2	243.0
Model / version selection guidance	8	5	8	2	160.0
Use-case-first onboarding redesign	8	9	8	4	144.0
In-dashboard usage analytics	7	7	8	4	98.0
Self-serve Growth / mid-tier SLA plan	6	9	7	5	75.6
AI onboarding assistant (chat-guided setup)	6	7	5	8	26.3
New docs visual theme / rebrand	3	2	9	3	18.0

11.3 Ranked priority order

1

Live credit / cost estimator QUICK WIN

Cheap, high-confidence, and it resolves a real trust gap at the widest point of the funnel.

2

Model / version selection guidance QUICK WIN

Another cheap fix that meaningfully reduces first-call confusion — a banner or tooltip, not a project.

3

Use-case-first onboarding redesign FLAGSHIP

Highest strategic impact at moderate effort. This is the Section 10 pick.

4

In-dashboard usage analytics

A solid mid-priority engineering investment that pays back in reduced support load as well as retention.

5

Self-serve Growth / mid-tier plan

Highest *business* impact of the set — it unlocks a whole revenue tier — but needs pricing, legal, and engineering alignment, which is why effort is scored 5.

6

AI onboarding assistant

Interesting but low confidence and high effort. Revisit once the items above have shipped and produced evidence.

7

Docs visual rebrand

A nice-to-have, correctly ranked last. Included deliberately: a matrix where every row scores well is a matrix that isn't doing any work.

WHY THE FLAGSHIP ISN'T THE #1 SCORE

The use-case-first onboarding redesign was selected for Section 10 rather than the top-scoring estimator because it best demonstrates end-to-end product and design thinking. In a real roadmap the two quick wins would ship first as NOW items — which is exactly how **Section 12** sequences them. Being explicit about that difference is the point: RICE ranks work, it doesn't choose what to showcase.