

Business Model

Who pays, what drives revenue and cost, and the uncomfortable arithmetic showing that self-serve developers are an ecosystem investment rather than a revenue engine.

MOTIONS	PRICING BASIS	FREE TIER	EST. SELF-SERVE ARPU
Two, running simultaneously	Consumption, in ₹	₹1,000, non-expiring	₹230–385 / month

DISCLAIMER

Sarvam AI is privately held and its internal unit economics are not public. This section uses **published pricing** to build a directional, illustrative model. None of these figures should be presented as confirmed company financials.

13.1 Who pays, and why

PAYER	WHY THEY PAY
Developers and startups	Pay-as-you-go API access — cheaper and faster than building Indic AI in-house
Enterprises	Contracted, compliant, scalable multilingual voice agents that cut call-centre cost
Government	Mission-funded or procurement-based sovereign AI deployment for citizen services

13.2 Published pricing snapshot

API	PRICE (₹, AS OF AUGUST 2026)
Chat completion (LLM)	₹4 / 1M input tokens · ₹2.5 / 1M cached input · ₹16 / 1M output tokens
Speech-to-text (Saaras)	₹30–45 per hour of audio, varying by feature such as diarisation
Text-to-speech (Bulbul)	₹15–30 per 10,000 characters
Translation	₹20–30 per 10,000 characters
Document parsing (Sarvam Vision)	₹0.5 per page, maximum 10 pages per job
Free tier	₹1,000 credits on signup, universal across APIs, never expire

13.3 What drives revenue and cost

Revenue drivers	UP	Cost drivers	DOWN
VOLUME	More tokens, minutes, characters, and pages processed — a classic consumption model	COMPUTE	Training, partly subsidised via IndiaAI Mission GPU capacity, plus recurring inference serving
CONTRACT	Enterprise deal value driven by number of languages, conversation volume, and deployment complexity (cloud vs. air-gapped)	FDE	Forward-deployed engineers embedded with enterprise customers — a higher-touch, higher-cost motion than pure self-serve SaaS
MIX	Shift toward enterprise, which carries far higher ARPU given banking, insurance, and government deal sizes	R&D	Training frontier-scale models (105B, and an announced 1T+ parameter model) is capital-intensive
		PEOPLE	Reported 226% year-on-year headcount growth adds significant operating cost as enterprise delivery scales

13.4 Illustrative unit economics — self-serve

SCENARIO INPUT	ASSUMPTION
Average developer monthly usage	2M chat tokens + 5 hours of speech-to-text + 50,000 characters of text-to-speech
Approximate monthly cost to developer	≈ ₹8 (chat) + ₹150–225 (STT) + ₹75–150 (TTS) = ₹230–385 / month
Implication	Self-serve revenue per developer is small in absolute rupee terms. Volume and enterprise mix — not per-developer ARPU — drive this business.

A directional back-of-envelope calculation from public price sheets. Not a claim about actual average usage or real revenue.

13.5 The tension worth naming

TWO BUSINESS MODELS, ONE COMPANY

Sarvam runs a low-touch, low-ARPU self-serve API business *and* a high-touch, high-ARPU enterprise and government business at the same time. The product implication — tying directly back to **Section 12** — is that investment in self-serve UX mainly buys ecosystem value, community, and developer-driven distribution, while the enterprise motion is what likely drives the majority of revenue today.

Why still invest in self-serve?

Developers become the startups that become the enterprises. Self-serve is the cheapest possible top of the enterprise funnel, and the community that forms around it is a distribution asset competitors cannot buy.

What this rules out

Any argument that self-serve improvements will move company revenue this quarter. The honest case for the Section 12 roadmap is ecosystem compounding and enterprise pipeline — not near-term ARPU.