SECTION 16 / 18

AI-NATIVE MEASUREMENT

AI-Specific Product Metrics

The metrics a standard SaaS dashboard would miss — quality, performance, economics, and adoption — and the argument that for Sarvam they must be tracked per language, not in aggregate.

FAMILIES	METRICS DEFINED	CRITICAL CUT	FEEDS
5	13	Per language, not aggregate	North Star & Section 17 risks

16.1 The framework

Quality — FAMILY 01

Q1	Task success rate
Q2	Accuracy — WER, BLEU
Q3	Hallucination / error rate

Performance — FAMILY 02

P1	Latency — p50 / p95
P2	Throughput at peak

Economics — FAMILY 03

E1	Cost per request
E2	Cost per active developer
E3	Gross margin per API product

Adoption — FAMILY 04

A1	Active developers and enterprises
A2	API calls per week, split by success and failure
A3	Activated accounts %

Retention — FAMILY 05

R1	Weekly Active Developers
R2	Returning API users %

16.2 Definitions and why each matters here

METRIC	DEFINITION	WHY IT MATTERS SPECIFICALLY FOR SARVAM
Task success rate	% of API calls or agent conversations that complete without error and satisfy user intent	The core input to the proposed North Star Metric in Section 6
Accuracy — WER, BLEU	Word Error Rate for speech-to-text; BLEU for translation quality	Sarvam's entire differentiation claim rests on being *more* accurate on Indic and code-mixed input than global competitors. This must be tracked per language, never only in aggregate.
Hallucination / error rate	% of chat or LLM responses containing fabricated or incorrect information	Highest-stakes for banking, insurance, and government. One bad hallucination in a loan-servicing agent is a compliance incident, not a bug report.
Latency — p50 / p95	Median and 95th-percentile response time	Voice agents fail as products when latency breaks the feel of real conversation. This is a UX metric disguised as an infrastructure metric.

METRIC	DEFINITION	WHY IT MATTERS SPECIFICALLY FOR SARVAM
Throughput	Maximum requests per second the system can serve reliably	Determines whether Sarvam can support a national-scale enterprise rollout without degradation
Cost per request	₹ cost to serve one API call, by product	Comes directly from published pricing in Section 13; needed to model gross margin per API line
Cost per developer	Total serving cost ÷ active users	The key to knowing whether ₹1,000 of free credit per signup remains sustainable at scale
Gross margin	(Revenue − serving cost) ÷ revenue, per API product	Chat, ASR, TTS, and OCR have very different compute intensity. A PM should never treat "Sarvam" as one margin profile.
Active developers	Distinct accounts making ≥1 successful call or agent conversation in a period	The core adoption signal; ladders into the metrics tree in Section 7
API calls per week	Raw volume, segmented by success and failure	A leading indicator of both adoption and infrastructure load
Activated accounts %	% of signups reaching a first successful call	The exact metric targeted by the Section 10 redesign and Section 14 Experiment 1
Weekly Active Developers	Developers making ≥1 successful call in a trailing 7-day window	The primary retention metric — the AI-era equivalent of DAU/WAU
Returning API users %	% of active developers or enterprises still active in the *following* period	The true stickiness signal, distinct from raw volume growth

16.3 The argument this section is really making

QUALITY METRICS ARE GROWTH METRICS

For an AI product, accuracy and hallucination rate are not engineering hygiene sitting beside the growth dashboard — they *are* growth metrics. A model that degrades on Tamil loses Tamil-speaking users before any signup number moves. Tracking quality per language is what makes that visible in time to act.

What aggregate metrics hide

A 4% aggregate error rate can comfortably conceal a 15% error rate in one language serving forty million people. For a company whose stated purpose is linguistic inclusion, that is not a rounding error — it is a failure of the mission, invisible on the dashboard.

What this enables downstream

Per-language tracking is also the mitigation for the bias risk named in **Section 17**. The metric and the safeguard are the same instrument, which is the cheapest kind of risk control there is.