


Risks

Seven AI-specific risks and six business risks, each with a mitigation direction — plotted so it is obvious which ones deserve engineering time this quarter and which only need watching.

AI RISKS 7	BUSINESS RISKS 6	HIGHEST SEVERITY Hallucination in regulated flows	MOST UNDER-DISCUSSED Enterprise concentration
---------------	---------------------	--	--

17.1 Risk matrix



POSITIONS REFLECT PRODUCT JUDGEMENT FROM PUBLIC INFORMATION, NOT MEASURED PROBABILITIES

17.2 AI-specific risks

RISK	DESCRIPTION	MITIGATION DIRECTION
Hallucination	LLM-based agents such as Samvaad and Indus may generate confidently wrong information — especially high-stakes in banking, insurance, and government contexts	Human-in-the-loop review for high-stakes flows; strict eval gating before feature launch; clear fallback-to-human paths
Privacy	Voice and document data from citizens and bank customers is highly sensitive	Sovereign, air-gapped deployment — already a core differentiator; strict compliance with the Digital Personal Data Protection Act
Bias	Models may perform unevenly across India's many languages and dialects, disadvantaging particular regions or communities	Per-language accuracy tracking (Section 16); prioritise under-served languages in eval sets and training data

RISK	DESCRIPTION	MITIGATION DIRECTION
Reliability	Outages or degraded quality directly break customer-facing banking and insurance flows	Defined SLAs — currently missing at the self-serve mid-tier per Section 9; uptime monitoring and graceful degradation
Cost	Serving frontier-scale models (105B, and the announced 1T+ parameter model) is expensive; free-credit generosity could pressure margins	Track cost-per-request and gross margin per API line (Section 16); tiered pricing that aligns cost with value delivered
Latency	Voice-agent UX depends on near-real-time response; larger models risk feeling sluggish	Model-size and latency trade-off guidance for developers (Section 15); Mixture-of-Experts architecture choices
Model drift	Model behaviour and quality can degrade or shift as usage patterns, data, or fine-tuning change over time	Continuous eval monitoring; versioned model releases with clear changelogs, already partially in place via the docs changelog

17.3 Business and market risks

RISK	DESCRIPTION
Competition from Krutrim and IndiaAI-backed peers	The closest strategic competitor is domestic, not global — a race for the same sovereign-AI narrative and the same government relationships, where being second is expensive
Hyperscaler bundling	AWS, Google, and Azure could bundle "good enough" Indic AI with infrastructure enterprises already buy, undercutting Sarvam on distribution rather than on quality
Dependence on government policy and funding	Significant tailwinds come from IndiaAI Mission compute subsidies. A policy shift could remove a meaningful cost advantage with little warning.
Enterprise concentration	High-value revenue is likely concentrated in a relatively small number of large enterprise and government contracts. Losing a major account could be material.
Talent and scaling risk	Reported 226% year-on-year headcount growth is fast. There is real execution risk in maintaining quality and culture while scaling forward-deployed engineering teams.
Self-serve monetisation risk	As shown in Section 13, self-serve ARPU appears low. If enterprise deal flow slows, the self-serve business alone would not be commercially sufficient.

THE RISK THAT CONSTRAINS THIS ROADMAP

Hallucination in regulated deployments is the one risk that can override product prioritisation. Any self-serve growth push proposed in **Section 12** must not dilute eval rigour on banking, insurance, and government surfaces — because a single compliance incident there costs more than every activation improvement in this study is worth.