SECTION 18 / 18

RECOMMENDATION

Executive Summary

The whole case study on one page — problem, evidence, insight, recommendation, metric, experiment, risk, and decision. Designed to be delivered aloud in under two minutes.

FORMAT	DELIVERY TIME	SYNTHESISES	DECISION
Interview cheat sheet	Under 2 minutes	Sections 1–17	Ship NOW items, greenlight redesign

18.1 The recommendation on one page

FIELD	CONTENT
Problem	Self-serve developers face high friction discovering *which* of Sarvam's six-plus APIs fits their use case, and then hit a "contact sales" wall when scaling past free-tier limits before they are enterprise-ready.
Evidence	Docs are organised by product rather than by use case (Section 9); free-credit value is never quantified at signup; no defined self-serve SLA tier exists between pay-as-you-go and full enterprise contracts.
Insight	Developers think in *jobs* — "I need a voice bot" — not in product catalogues. The funnel currently forces them to do the API-matching work themselves before they have experienced any value at all.
Recommendation	Ship a use-case-first onboarding redesign (Section 10) alongside two quick NOW wins — a live credit estimator and model-selection guidance (Section 12) — then introduce a self-serve "Growth" mid-tier to close the scaling gap.
Expected impact	Meaningful lift in first-successful-call rate and 7-day developer activation, plus incremental self-serve revenue from developers who currently churn at the "contact sales" wall.
Metric	Proposed North Star — **Weekly Successful AI Tasks Completed** (Section 6), moved via the Activation and Revenue-tier branches of the metrics tree (Section 7).
Experiment	A/B test use-case-first onboarding against the current docs-first flow, with first-successful-call rate as the primary metric and failed-call plus support-ticket rate as the guardrail (Section 14).
Risk	Hallucination and reliability risk remain highest in enterprise and government voice-agent deployments. Any self-serve growth push must not dilute eval rigour on those high-stakes surfaces (Section 17).
Decision	Prioritise the NOW items immediately — low effort, high confidence. Greenlight the onboarding redesign as the flagship NEXT initiative. Hold the Growth plan and the AI onboarding assistant as LATER bets pending validation.

18.2 The argument in four beats

1. **Sarvam's moat is real but narrow**
   Indic-language depth and sovereign deployment are genuinely defensible. General frontier capability is not, and the roadmap should stop trying to compete there.

2. **The self-serve funnel leaks at a fixable place**
   The gap between how developers state their job and how the product presents itself is an information-architecture problem, not a model problem — which makes it cheap to fix.

3. **The business case is ecosystem, not ARPU**
   Self-serve revenue per developer is small. The honest justification for this work is enterprise pipeline and community distribution, and saying so up front makes the recommendation more credible, not less.

4. **Quality is the constraint on growth**
   Per-language accuracy and hallucination rate cap how fast any of this can move. They belong on the growth dashboard, not beside it.

18.3 Why Sarvam is a strong PM case study

Multi-sided go-to-market

Developer self-serve, enterprise sales, and government procurement inside a single product organisation — three motions that pull the roadmap in different directions and cannot all be satisfied at once.

AI-specific quality metrics

Hallucination rate and per-language accuracy sitting alongside standard growth metrics, with the argument that in an AI product these are the same category of number.

A real moat weighed against real gaps

Sovereignty and Indic depth on one side, documentation maturity and general capability on the other. Holding both is what makes the analysis balanced rather than promotional.

Business-model bifurcation

Low-ARPU self-serve against high-ARPU enterprise — a pattern that recurs across most real B2B and B2B2C AI platforms, which makes the reasoning here transferable.

> **IF ASKED ONLY ONE QUESTION**
> "What would you ship first, and why?" — The live credit estimator. It is two weeks of work, it resolves a trust gap affecting every single signup, and it produces the fastest evidence about whether the deeper onboarding thesis in Section 10 is correct before anyone commits a quarter to building it.